

$$F = G \frac{m_1 m_2}{d^2}$$

Searching and heuristics

$$-E + V = 2$$

$$E = mc^2$$

$$ds \geq 0$$

$$\frac{\partial^2 u}{\partial t^2} = c^2 \frac{\partial^2 u}{\partial x^2}$$

March 2025

$$\frac{df}{dt} = \lim_{h \rightarrow 0} \frac{f(t+h) - f(t)}{h}$$

Topics to be discussed

Part I. Introduction

1. Introduction to Artificial intelligence

Part II. Search and optimisation.

2. Searching and heuristics
3. Search - optimisation
4. Two-player deterministic games
5. Evolutionary and genetic algorithms

Part III. Machine learning and data analysis.

6. Regression, classification and clustering (Part I & II)
8. Artificial neural networks
9. Bayesian models
10. Reinforcement Learning

Part IV. Logic, Inference, Knowledge Representation

11. Propositional logic and predicate logic
12. Knowledge Representation

Part V. AI in Action: Language, Vision

13. AI in Natural language processing
14. AI in Vision and Perception

Part VI. Summary

15. AI engineering, Explainable AI, Ethics

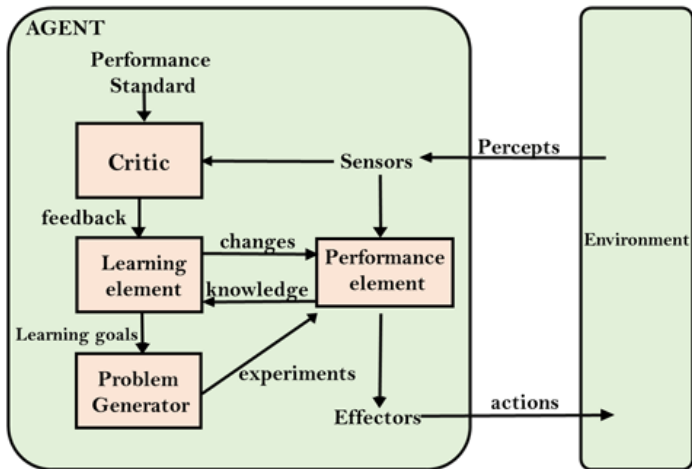
Overview

- Introduction
- Types of basic search methods in AI
 - Uninformed search methods
 - Informed search methods
- Comparing uninformed and informed search methods
- Applications of search methods in AI

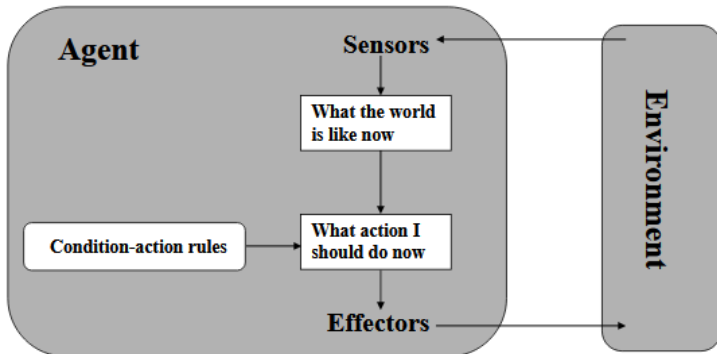
Two views of AI

- **AI agents:** how can we create intelligence?
 - The inspiration comes from the types of capabilities that humans possess.
- **AI tools:** how can we benefit society?
 - Need comes from existing problems in the world, and techniques developed by the AI community happen to be useful for that.

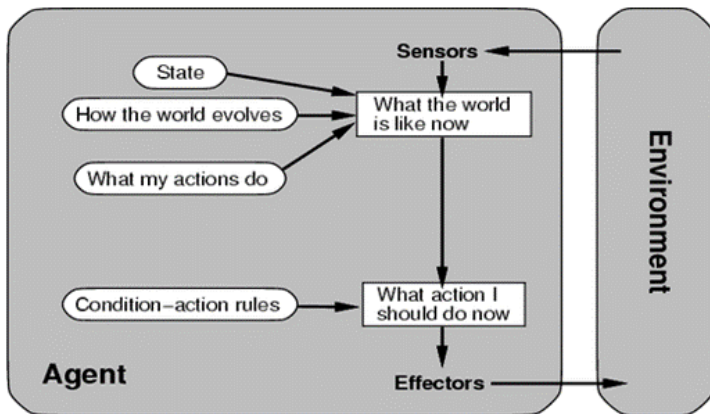
Intelligent Agents in AI



A simple Reflex Agent



Agent that keeps track of the world



What are search methods in AI?

- Search methods are techniques used by AI systems to find solutions to problems by searching through a set of possible solutions.
- These methods involve generating a set of candidate solutions and evaluating them based on some objective criteria.
- Different types of basic search methods exist, from uninformed methods like depth-first search and breadth-first search to informed methods like A* search and best-first search.


Examples of classical” applications of search in AI

- In games such as chess or checkers, AI programs use search methods to explore different possible moves and their outcomes to find the best move.
- In natural language processing, search methods can be used to find the most likely interpretations of a given sentence by searching through possible syntax trees or semantic representations.
- In robotics, search methods can be used to plan the most efficient path for a robot to navigate through an environment by searching through a set of possible paths.
- In computer vision, search methods can be used in applications such as face recognition by searching through a database of images to find a match.



Example: Toy problems

- N queens puzzle: how to place N queens on a $N \times N$ board so that none is attacked by any other
- Two containers have a capacity of 5 and 7 liters. How, by filling them to the full, emptying and pouring water between them, measure 4 liters?
- Traveling salesman problem

8 puzzle problem



Graphs and networks



State graph definitions (1)

- **Graph:** A set of vertices V and a set of edges E , where each edge is a pair (u, v) representing a connection between vertices u and v .
- **Node/Vertex:** A point in the graph.
- **Edge:** A connection between two vertices u and v .
- **Directed/Undirected Graph:** A directed graph has edges with a direction, while an undirected graph has edges without a direction.
- **Adjacency List:** A data structure used to represent a graph as a list of vertices, where each vertex has a list of its adjacent vertices.
- **Adjacency Matrix:** A two-dimensional matrix used to represent a graph, where the rows and columns correspond to the vertices, and the entries represent the existence of an edge between the corresponding vertices.
- **Traversal:** Visiting all the nodes/vertices of a graph in a systematic order.


State graph definitions (2)

- **Stack:** A data structure that follows the Last-In-First-Out (LIFO) principle, i.e., the last element added to the stack is the first one to be removed.
- **Queue:** A data structure that follows the First-In-First-Out (FIFO) principle, i.e., the first element added to the queue is the first one to be removed.
- **Visited Array/Set:** A data structure that keeps track of the nodes/vertices that have already been visited by an algorithm.
- **Path:** A sequence of nodes/vertices that represents a route between two nodes/vertices in a graph.
- **Cost/Distance:** A numerical value associated with an edge that represents the "cost" or "distance" of traversing that edge.
- **Search Space:** The set of all possible states/nodes that can be explored by a search algorithm.
- **Heuristic Function:** A function that estimates the "cost" or "distance" of reaching the goal node from a given node, and is used to guide a search algorithm towards the goal node.

Two basic graph search algorithms

- **BFS (Breadth-First Search):** A graph traversal algorithm that visits all the nodes/vertices of a graph in breadth-first order, i.e., it visits all the nodes at the current depth/level before moving on to the nodes at the next depth/level.
- **DFS (Depth-First Search):** A graph traversal algorithm that visits all the nodes/vertices of a graph in depth-first order, i.e., it explores as far as possible along each branch before backtracking.

Depth-First Search, FS



Depth-First Search, DFS

Algorithm 1: DFS Algorithm

Result: Visited vertices in DFS order

initialize stack with starting vertex

initialize visited set as empty

initialize result list as empty

while *stack is not empty* **do**

 pop vertex v from stack

if *vertex v has not been visited* **then**

 mark vertex v as visited

 add vertex v to result list

foreach *neighbor u of vertex v in reverse order* **do**

if *vertex u has not been visited* **then**

 push vertex u onto stack

 add dotted arrow from v to u


end

end

end

end

Breadth-First Search, BFS



Breadth-First Search, BFS

Algorithm 2: Breadth-First Search

Input: Graph $G = (V, E)$ and starting vertex s

Output: Visited vertices

foreach vertex $v \in V$ **do**

$visited[v] \leftarrow false$;

end

$queue \leftarrow [s]$;

$visited[s] \leftarrow true$;

while $queue \neq []$ **do**

$u \leftarrow queue.pop(0)$;

foreach vertex v adjacent to u **do**

if $visited[v] = false$ **then**

$visited[v] \leftarrow true$;

$queue.append(v)$;

end


end

end

BFS vs. DFS - properties summary

- BFS finds a shortest sequence, it needs less time but more memory
- DFS needs less memory, but is not guaranteed to find a shortest sequence

Iterative Deepening in Depth-First Search



Iterative Deepening Depth-First Search, IDDFS

Algorithm 3: Iterative Deepening Depth-First Search (IDDFS)

Data: Graph G , Start vertex s

Result: The sequence of visited vertices during the search

Function $\text{DFS}(G, s, \text{depth})$:

if $\text{depth} \leq 0$ **then**

return

 Mark s as visited;

forall neighbors n of s **do**

if n is not visited **then**

$\text{DFS}(G, n, \text{depth}-1)$;

for $\text{depth} = 1$ **to** ∞ **do**

 Run DFS on G with starting vertex s and maximum depth limit depth ;

if all vertices are visited **then**

return visited sequence;

Bidirectional search

- **Bidirectional search** is a graph search algorithm that searches two ends of the graph simultaneously, starting from the initial and goal nodes.
- In each step of the search, it expands the node that has not been expanded before and has the lowest cost.
- Once both searches meet at a common node, the search is complete and returns the path from the start to the goal node.
- Bidirectional search is particularly useful when the search space is large, and the branching factor is high, as it can significantly reduce the number of nodes that need to be explored.
- It can also be more efficient when the graph is structured, such as in a tree or a grid, and the start and goal nodes are close to each other.
- However, bidirectional search requires the ability to efficiently check if a node has been visited, and it also requires a way to efficiently check if a node is a common node between the two searches.

Bidirectional Search Algorithm

Algorithm 4: Bidirectional Search Algorithm

Input: start node s and goal node g

Output: shortest path from s to g

$forwardFrontier \leftarrow \{s\};$

$backwardFrontier \leftarrow \{g\};$

$forwardExplored \leftarrow \{\};$

$backwardExplored \leftarrow \{\};$

while $forwardFrontier \neq \emptyset$ and $backwardFrontier \neq \emptyset$ **do**

if $|forwardFrontier| \leq |backwardFrontier|$ **then**

$current \leftarrow$ a node from $forwardFrontier$;

 remove $current$ from $forwardFrontier$;

if $current \in backwardFrontier$ or $current \in backwardExplored$ **then**

return path from s to $current$ and path from g to $current$;

end

 add $current$ to $forwardExplored$;

for each neighbor n of $current$ not in $forwardExplored$ **do**

 add n to $forwardFrontier$;

end

end

else

$current \leftarrow$ a node from $backwardFrontier$;

 remove $current$ from $backwardFrontier$;

if $current \in forwardFrontier$ or $current \in forwardExplored$ **then**

return path from s to $current$ and path from g to $current$;

end

 add $current$ to $backwardExplored$;

for each neighbor n of $current$ not in $backwardExplored$ **do**

 add n to $backwardFrontier$;

end

end


end

return no path found;

Finding the Shortest Path

- Suppose we have a network of cities and want to find the shortest path between two of them.
- We can represent the network as a graph, where each city is a node and the distance between them is the cost of the edge connecting them.
- To find the shortest path, we can use a search algorithm that explores the graph and keeps track of the current best path.
- One such algorithm is Dijkstra's algorithm, which starts at the source city and repeatedly selects the node with the lowest distance until the destination city is reached.
- Alternatively, we can use a heuristic search algorithm like A^* that incorporates information about the estimated remaining distance to the goal city to guide the search.

Uniform cost search – an example



Node	Visited?	Total Cost
A	Yes	0
G	Yes	20
D	Yes	38
B	Yes	41
E	Yes	45
I	Yes	52
J	Yes	61
C	Yes	85
F	Yes	88
H	Yes	100
K	Yes	123

Uniform cost search

Algorithm 5: Uniform Cost Algorithm

Input: start node s , goal node g

Data: edge cost function c , set of nodes V , set of edges E

Output: shortest path from s to g

```
foreach node  $v$  in  $V$  do
```

$$d[v] \leftarrow \infty ;$$

```
// initialize all distances as infinity
```

end

$$d[s] \leftarrow 0;$$

```
// distance from start to start is 0
```

```
frontier ← PriorityQueue();
```

```
// initialize the priority queue
```

```
frontier.put(s, 0) ;
```

```
// put the start node in the queue with priority 0
```

```
while frontier is not empty do
```

```
u ← frontier.get() ;
```

```
// get the node with the lowest cost so far
```

if $u = g$ then

```

| return construct_path(s, g) ;

```

```
// goal reached, return the path
```

end

```
foreach neighbor  $v$  of  $u$  do
```

$$cost \leftarrow c(u, v);$$

```
// get the cost of the edge from u to v
```

```
new_distance  $\leftarrow d[u] + cost$  ;
```

```
// calculate the new distance to v via u
```

```

if new_distance < d[v] then

```

$$d[v] \leftarrow new_distance ;$$

```
// update the distance to v
```

```
frontier.put(v, d[v]) :
```

```
// add v to the queue with updated priority
```

$$prev[v] \leftarrow u;$$

```
// update the parent of v to u
```

end

end

end

```

    return None :

```

```
// goal not reachable from start
```

Greedy Best First Algorithm

- Greedy Best-First Search is an informed search algorithm that uses a heuristic function $h(n)$ to guide the search towards the goal state.
- At each step, the algorithm expands the node with the lowest heuristic value, i.e., the node that is closest to the goal according to the heuristic function.
- The algorithm maintains a priority queue called the frontier, which contains the nodes to be expanded in the search tree.
- The heuristic function must be admissible, i.e., it should never overestimate the true cost of reaching the goal state.
- Greedy Best-First Search is not guaranteed to find the optimal solution, but it can be very efficient if the heuristic function is well-designed.

Greedy Best First Algorithm

Algorithm 6: Greedy Best-First Search

Input: initial state s_0 , heuristic function h

Output: goal state, or failure

$frontier \leftarrow \{s_0\}$

$explored \leftarrow \emptyset$

while $frontier \neq \emptyset$ **do**

$s \leftarrow$ state in $frontier$ with lowest $h(s)$

if s is goal state **then**

return s

 Move s from $frontier$ to $explored$

for each successor s' of s **do**

if $s' \notin frontier$ or $h(s') < h(s)$ **then**

 Add s' to $frontier$

return failure

A* algorithm

- A* pathfinding algorithm that finds the shortest path between two nodes in a weighted graph
- Works by combining the cost of the path from the start node to the current node (g-value) with the estimated cost of the path from the current node to the goal node (h-value)
- The f-value of a node is the sum of the g-value and the h-value:
$$f(n) = g(n) + h(n)$$
- Keeps track of the f-values of all nodes in a priority queue, sorted by lowest f-value
- Expands the node with the lowest f-value, updating the f, g, and h values of its neighbors as needed
- Stops when the goal node is reached, or all paths have been explored

Algorithm 7: A* Algorithm

Input: Start node s , goal node g , heuristic function h , cost function c

Output: Shortest path from s to g

$frontier \leftarrow \{s\}$

$came_from \leftarrow \{\}$

$cost_so_far[s] \leftarrow 0$

while $frontier$ is not empty **do**

$current \leftarrow$ node in $frontier$ with lowest f value

if $current = g$ **then**

return $reconstruct_path(came_from, s, g)$

 remove $current$ from $frontier$

foreach neighbor $next$ of $current$ **do**

$new_cost \leftarrow cost_so_far[current] + c(current, next)$

if $next$ not in $cost_so_far$ or $new_cost < cost_so_far[next]$ **then**

$cost_so_far[next] \leftarrow new_cost$


$priority \leftarrow new_cost + h(next, g)$

 add $next$ to $frontier$ with priority $priority$

$came_from[next] \leftarrow current$

return $null$

A* algorithm – an example



The heuristic h function

- Relatively simple to evaluate
- Can be a solution to a simplified problem
 - By ignoring certain constraints
 - It is well if the simplified problem has an analytical solution
 - It can be evaluated through on-line solving of the simplified problem

IDA* (Iterative Deepening A*)

- An informed search algorithm that finds the shortest path between two nodes in a weighted graph
- Combines the advantages of A* and iterative deepening depth-first search
- Works by searching the graph repeatedly with increasing depth limits until a solution is found
- Uses a heuristic function to estimate the remaining cost to reach the goal node, represented by $h(n)$
- The cost of the path from the start node to a node n is represented by $g(n)$
- The total estimated cost of the path through node n to the goal node is represented by $f(n) = g(n) + h(n)$
- At each iteration, nodes are expanded in increasing order of their f -values until the goal node is reached
- If the goal node is not reached within the depth limit, the search is repeated with a larger depth limit

IDA* (Iterative Deepening A*)

Algorithm 8: Iterative Deepening A* Algorithm

Input : start node s , goal node t

Output: shortest path from s to t

foreach node n **do**

$g(n) \leftarrow \infty$

$f(n) \leftarrow \infty$

$closed(n) \leftarrow \text{false}$

$g(s) \leftarrow 0$

$f(s) \leftarrow h(s)$

for $depth \leftarrow 0$ **to** ∞ **do**

$open \leftarrow$ priority queue with s as the only element

while $open$ is not empty **do**

$n \leftarrow$ node with minimum $f(n)$ value in $open$

if n is the goal node **then**

return the path from s to n

$closed(n) \leftarrow \text{true}$

foreach successor s of n **do**

if s is not in $closed$ or $g(s) > g(n) + w(n, s)$ **then**

$g(s) \leftarrow g(n) + w(n, s)$

$f(s) \leftarrow g(s) + h(s)$

if $depth > 0$ **then**

 add s to $open$

return failure

Beam Search

- Beam search is a heuristic search algorithm that explores a graph by expanding the most promising nodes at each level of the search.
- **Strengths:**
 - Beam search is an improvement over simple best-first search, as it uses a beam width parameter to reduce the search space and avoid getting stuck in local optima.
 - It can be easily parallelized by processing different branches of the search space concurrently.
 - It can handle large search spaces and can be effective for finding approximate solutions to problems where the exact solution is not required.
- **Weaknesses:**
 - Beam search is not guaranteed to find the optimal solution, as it may discard the correct path early in the search process.
 - The performance of beam search depends heavily on the choice of the beam width parameter, which can be difficult to determine in advance.
 - Beam search may get stuck in local optima if the search space has many similar promising nodes, and the beam width is too narrow.

Algorithm 9: Beam search algorithm

Input: Starting node s , beam width k , goal node g

Output: Shortest path from s to g

$B \leftarrow \{s\};$

while g not in B **do**

$B' \leftarrow \emptyset;$

foreach node n in B **do**

foreach successor s of n **do**

$B' \leftarrow B' \cup \{s\};$

end

end

 Sort nodes in B' based on cost from s + heuristic estimate to g ;

$B \leftarrow$ the k best nodes in B' ;

end

return shortest path from s to g found in B ;

Algorithm Comparison on time, memory, completeness and optimality

Algorithm	Time	Space	Complete	Optimal
Breadth First	$O(b^d)$	$O(b^d)$	Yes	Yes
Uniform Cost Search	$O(b^{1+\lceil \log(C^*/\epsilon) \rceil})$	$O(b^{1+\lceil \log(C^*/\epsilon) \rceil})$	Yes	Yes
Depth First	$O(b^m)$	$O(bm)$	No	No
Depth Limited	$O(b^l)$	$O(bl)$	No	No
Iterative Deepening	$O(b^d)$	$O(bd)$	Yes	Yes
Bidirectional	$O(b^{d/2})$	$O(b^{d/2})$	Yes	Yes

Comparison and recommendations

Algorithm	Strengths	Weaknesses
Best First	<ul style="list-style-type: none">• Can be implemented efficiently• Useful when optimality is not critical• Useful when heuristic function h is accurate	<ul style="list-style-type: none">• Does not guarantee optimal solution• May not be suitable for problems with large state spaces
A*	<ul style="list-style-type: none">• Guaranteed to find optimal solution, provided heuristic function h is admissible and consistent• Critical when optimality is important	<ul style="list-style-type: none">• Can be memory-intensive• May not be suitable for problems with large state spaces
IDA*	<ul style="list-style-type: none">• Limits amount of memory used by expanding nodes incrementally• Can be used even when memory requirements exceed available resources	<ul style="list-style-type: none">• Can be less efficient than A*• May require a large number of iterations to find optimal solution

Overview

- Introduction
- Types of basic search methods in AI
 - Uninformed search methods
 - Informed search methods
- Comparing uninformed and informed search methods
- Applications of search methods in AI

- ① S. J. Russell, P. Norvig, "Artificial Intelligence: A Modern Approach", Financial Times Prentice Hall, 2019.
- ② M. Flasiński, "Introduction to Artificial Intelligence", Springer Verlag, 2016
- ③ M. Muraszkiewicz, R. Nowak (ed.), "Sztuczna Inteligencja dla inżynierów", Oficyna Wydawnicza PW, 2022
- ④ J. Prateek, "Artificial Intelligence with Python", Packt 2017